

ROART: Range-query Optimized Persistent ART

Shaonan Ma¹, Kang Chen¹, Shimin Chen², Mengxing Liu¹, Jianglang Zhu¹, Hongbo Kang¹,
and Yongwei Wu¹

¹Tsinghua University

²SKL of Computer Architecture, ICT, CAS, and University of Chinese Academy of Sciences

FAST 2021

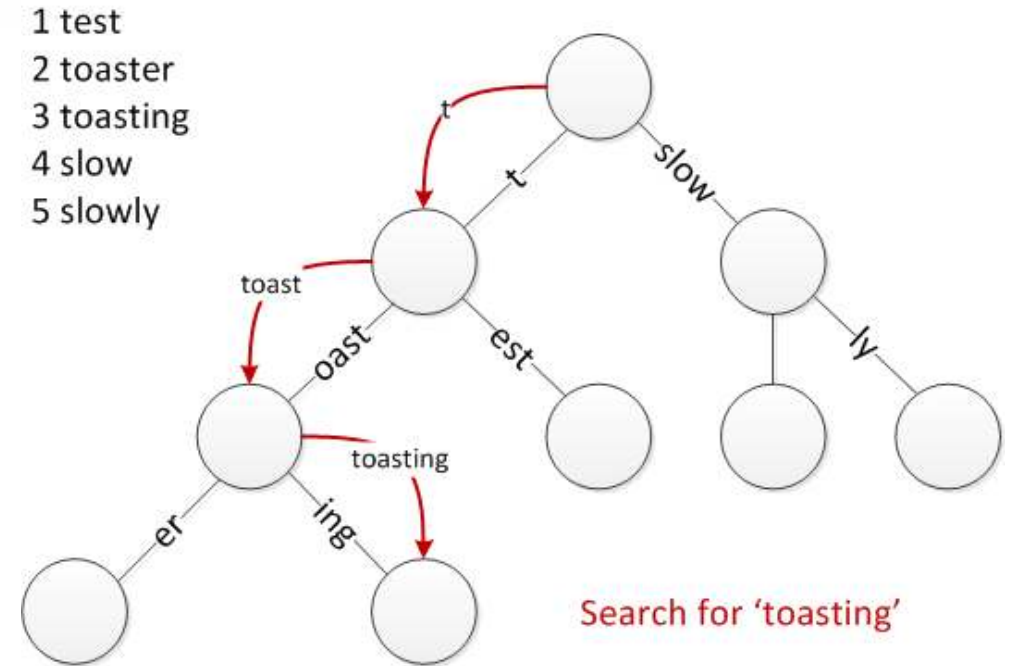
Background

➤ Radix Tree

- A space-optimized trie (prefix tree)
- Each internal node has **at most r** children ($r = 2^n, n \geq 1$)
- Edge: single element or sequences of elements

➤ Application


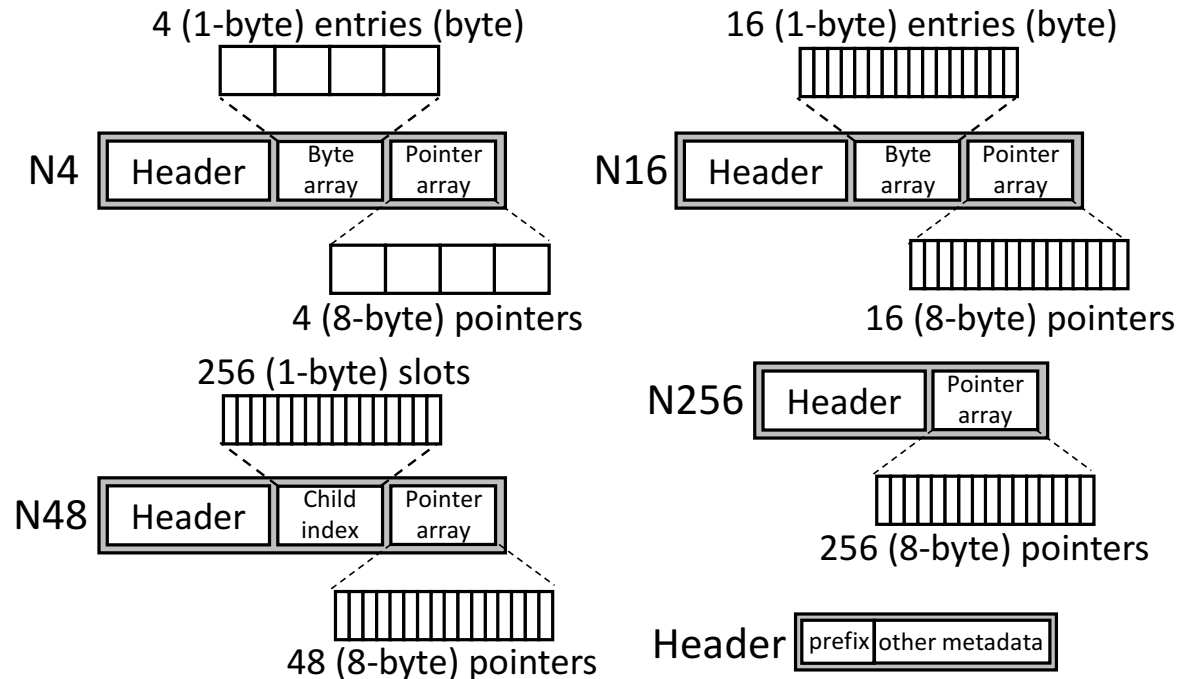
- Efficient for small sets or for sets of strings that share long prefixes
- URL matching (search engine)
- IP routing



ART

➤ ART: Adaptive Radix Tree

- Integrate **adaptive node sizes** to the radix tree
- Grow while adding new entries
- Better space of use without reducing speed compared with the radix tree




partial keys 0, 2, 3, and 255 are mapped to the subtrees a, b, c, and d

ART

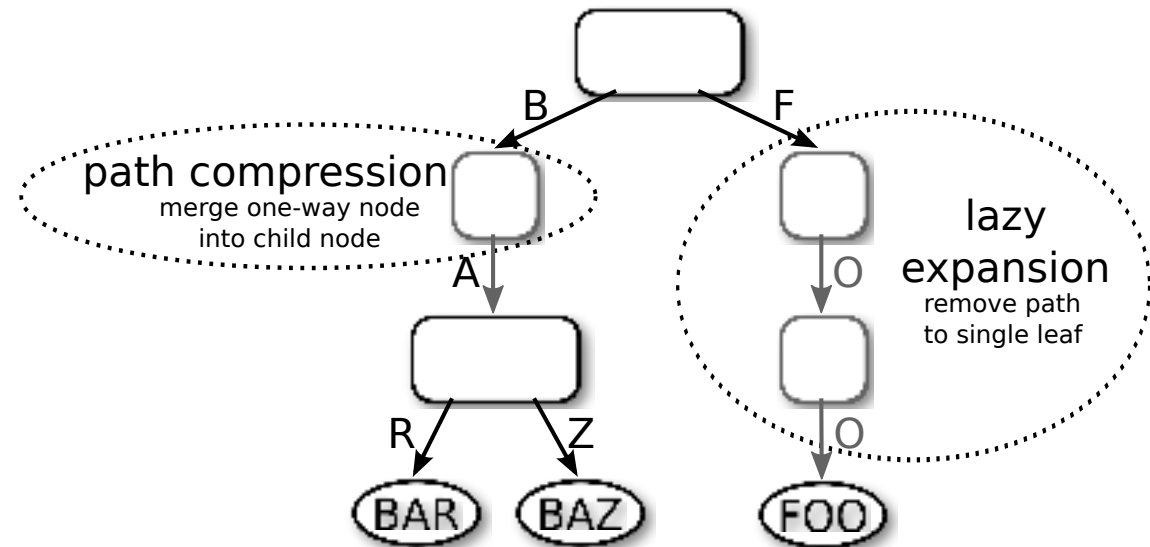
➤ Techniques in ART

- Path compression
- Lazy expansion

} Reduce height & Space



Adaptively sized nodes in **ART**



Path Compression: B-A-R → BA-R

Lazy expansion: F-O-O → F-OO


Persistent Memory

➤ Features

- Non-Volatile Memory
- Load/Store Instructions

➤ Advantage

- Lower access latency than SSDs
- Cheaper than DRAM
- Byte-addressable && Memory-like access




➤ Commercial persistent memory was available in 2019.04


Motivation I: Functionality

➤ Range Queries Support

- Requirement of most KV stores (Redis, Memcached) and databases (MySQL, PostgreSQL)



Pointer chasing in B⁺-Tree



- Keys can be in leaf node or non-leaf node → Pointer chasing
- String Comparison


Motivation I: Functionality

➤ Variable-sized Keys Performance

- **1.8-3.9X** Performance degradation in B⁺-Tree

➤ Existing Solutions

- Fixed-sized B⁺-Tree → **Optimize 8-byte keys only**
- Variable-sized B⁺-Tree (store the addresses of keys in indexes) → **Incur more pointer chasing**



FAST&FAIR (B⁺-Tree), 8-byte keys, 4-threads


Performance Degradation:

pointer chasing && string comparison during traversal

Motivation II: Performance

➤ Poor write scalability

- Reduce persistence overhead → crucial for performance




- Optane-NI: single persistent memory
- Optane: 6 interleaved persistent memory
- All threads use a 256 B access size

Persistent Memory has the **poor write scalability**
→ Avoid flush and memory fence as much as possible

Motivation III: Correctness

➤ Memory Safety

- **Volatile Allocator** → Memory Leaks
- **Persistent Allocator**
 - Logging-based
 - Post-crash GC




➤ Logging-based

- Constant time recovery
- Slow allocation/deallocation

➤ Post-crash GC

- Fast allocation/deallocation
- Slow recovery



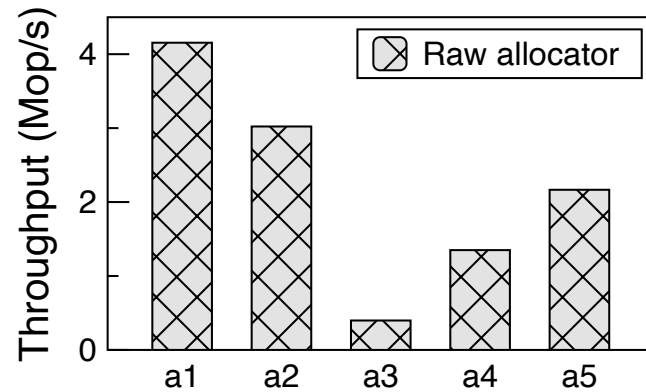
Makalu with post-crash GC

Motivation III: Correctness

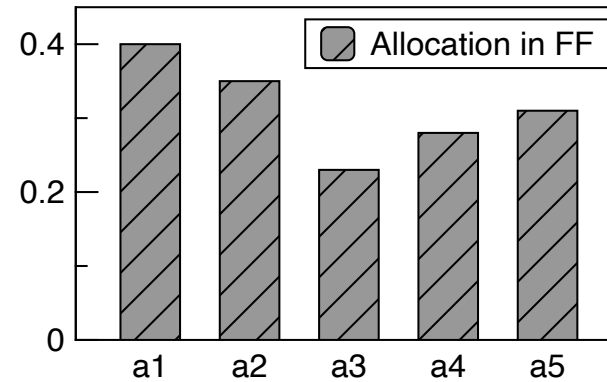
➤ Allocator Performance

- a1: **malloc**, the standard volatile allocator for DRAM
- a2: **libvmmalloc**, volatile allocator based on jemalloc
- a3: **PMDK**, logging-based persistent allocator
- a4: **nvm_malloc**, logging-based persistent allocator
- a5: **Makalu**, persistent allocator with post-crash GC

Allocate 64-byte chunk, then write and persist them in a single thread



(a) Performance of raw allocators



(b) Allocators in FAST&FAIR

(a) **Makalu** is **50%** and **28%** slower than **malloc** and **libvmmalloc**

(b) Gaps between different cases are narrowed → **Tree's traversal overhead**

Motivation

➤ Functionality

- Range queries support
- Reduce pointer chasing

→ Leaf compaction based on **ART**

➤ Performance

- Reduce PM writes/flush

→ Entry compression && Minimal persistence

➤ Correctness

- Fast allocation
- Prevent memory leaks
- Fast recovery

→ **DCMM** with instant restart

ROART

➤ **ROART**, a **R**ange **O**ptimized based on **A**daptative **R**adix **T**ree

- Reduce pointer chasing
- Reduce persistent overhead
- Tend to lower the height of tree

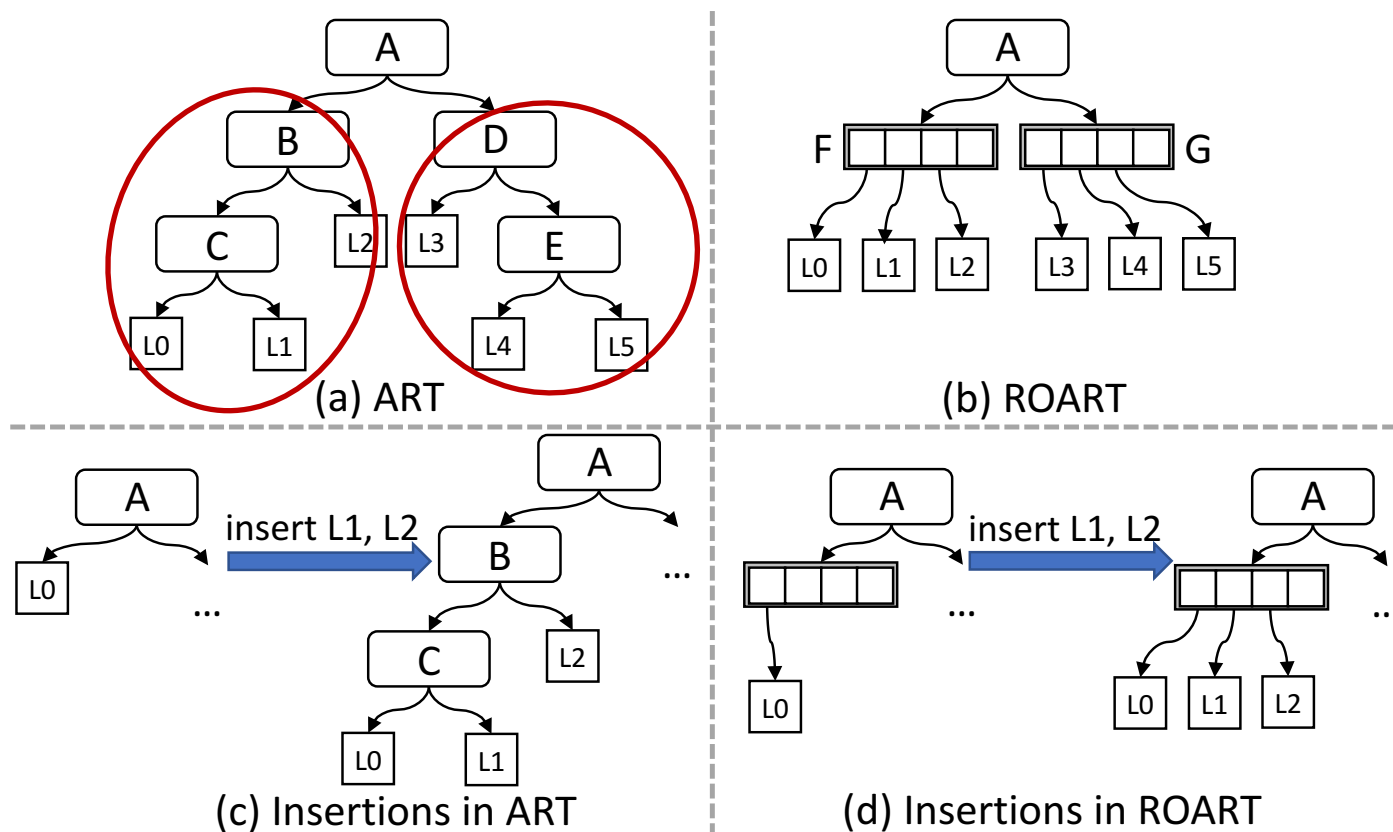
➤ Techniques

- Leaf compaction
- Entry Compression
- Selective Metadata Persistence
- Minimally Ordered Split
- Delayed Check Memory Management(**DCMM**)

Leaf Compaction

➤ Leaf compaction

- Compact the pointers of leaf nodes into a leaf array



If a subtree of the radix tree has $\leq m$ leaf nodes, the subtree is compacted into a leaf array ($m = 4$ in this case)

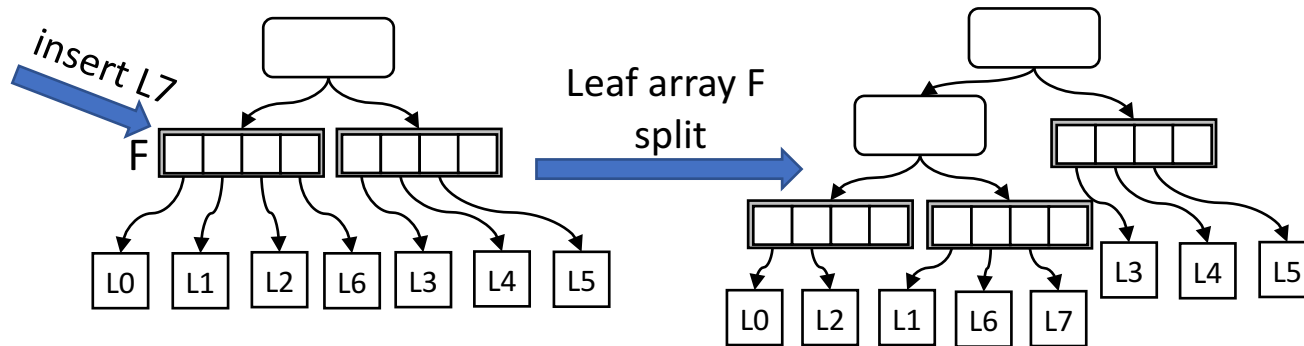
- **ROART** embeds **fingerprints** (hash) of key in leaf array to avoid comparison for most unnecessary cases

- ✓ Pointer chasing ↓
- ✓ Tree height ↓


Leaf Split

➤ Leaf Split according to first non-common bytes

- Split after the first non-common bytes
- Split cost is high but rare



Insert L7; L1, L6 and L7 share the same prefix



Example of leaf array split

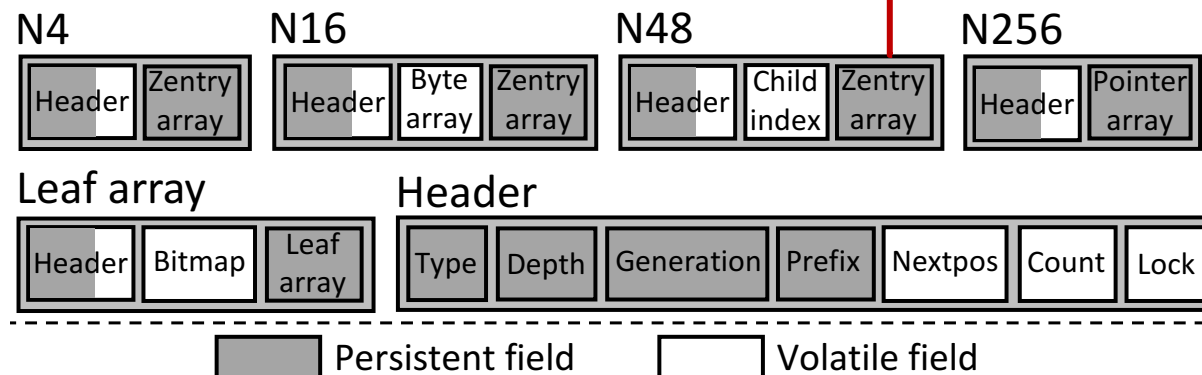
Optimize Persistent Overhead

➤ Entry Compression

- Compact key byte into the pointer in N4, N16, N48

➤ Selective Metadata Persistence

- Do not persist data which can be computed by scanning Zentry or pointer array
- Generation is used for lazy recovery (Global generation number vs Node generation number)
 - Global generation number += 1 when restarted
 - Restore metadata on demand



Per element in Zentry array:

empty: 8-bit	key: 8-bit	pointer: 48-bit
--------------	------------	-----------------

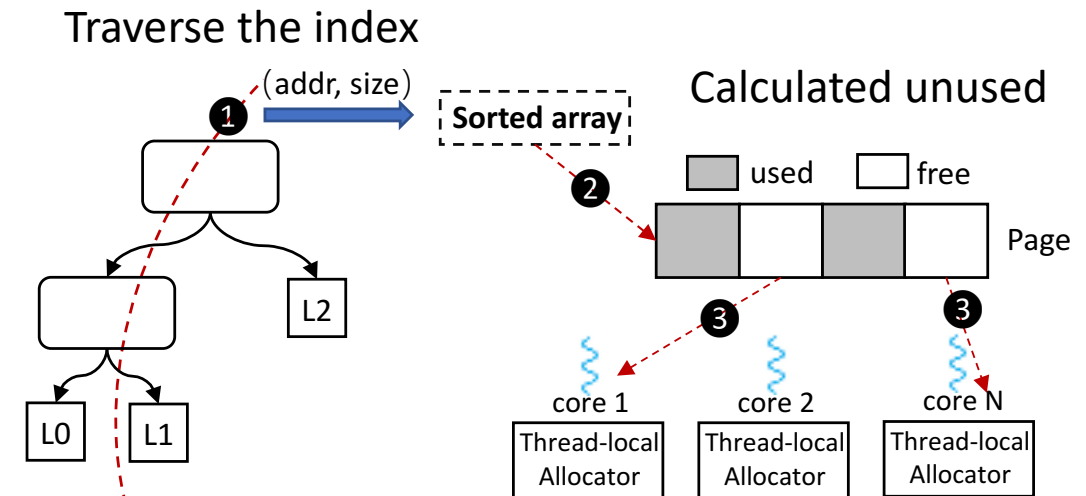
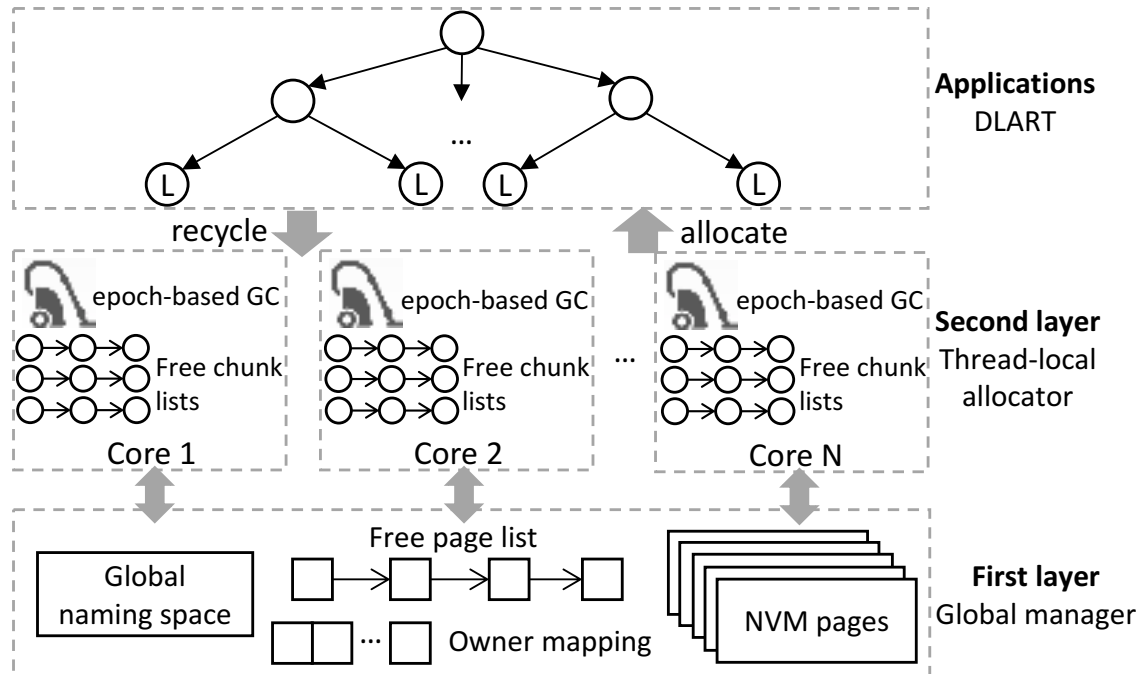
Additional optimization: Minimally Ordered Split

→ Concurrent split operations & Reduce fence

Node structures in **ROART**

Allocator: DCMM

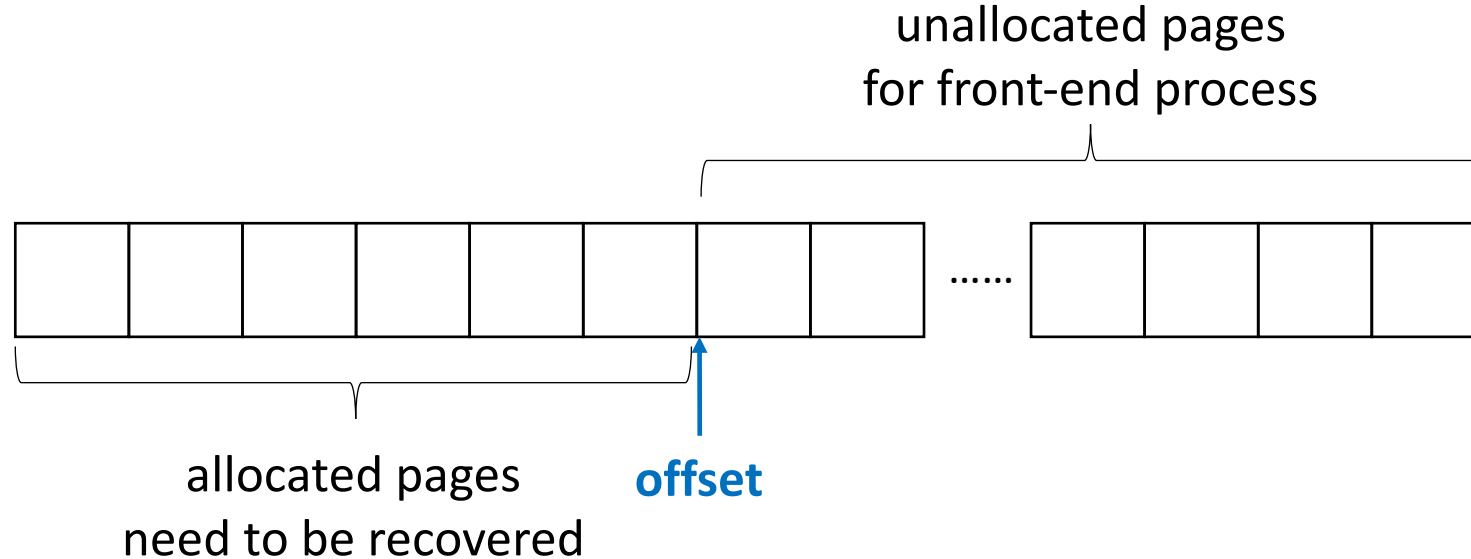
- **Delayed Check Memory Management (DCMM) based on post-crash GC**
 - Global naming space: Contain the roots of indexes and an **offset** field indicating the offset of last allocated page
 - Obtain offset & increase **offset** (persistence required) to request a new page



owner_mapping: a map between each page and its owner thread, [0, offset)

Instant Restart

- Allocate immediately new pages after offset **without waiting** for other metadata recovery to complete
- Additional: Lazy recovery
 - Global generation number (± 1 when restarted) \leftrightarrow Node generation number



Start from offset when recovery

Experimental Setup

➤ Evaluation Setup

- Dell PowerEdge R740 server
- Four Intel(R) Xeon(R) Gold 5220 processors
- 6×128GB Optane DC PMM per socket
- 32KB L1-cache, 1MB L2-cache, and 25MB L3-cache
- Persistent memory is managed by a DAX file system

➤ Comparison


- **P-ART** [SOSP'19]
 - **PMwCAS-ART**(based on PMwCAS [ICDE'18])
 - **FAST&FAIR** [FAST'18]
 - **Lock-free SkipList** [ATC'18]
 - **BzTree** [VLDB'18]
- Modify **P-ART**, **FAST&FAIR** and **SkipList** with **DCMM** allocator for a fair comparison.
 - **PMwCAS** and **BzTree** use their own persistent allocators based on **PMDK**

Overall Performance

➤ Methodology

- 4 threads to generate randomly
- **Keys:** randomly generated with sizes: 4-128 bytes
- **Values:** fixed as 8 bytes

- FAST&FAIR: binary search X
- SkipList: cache locality X



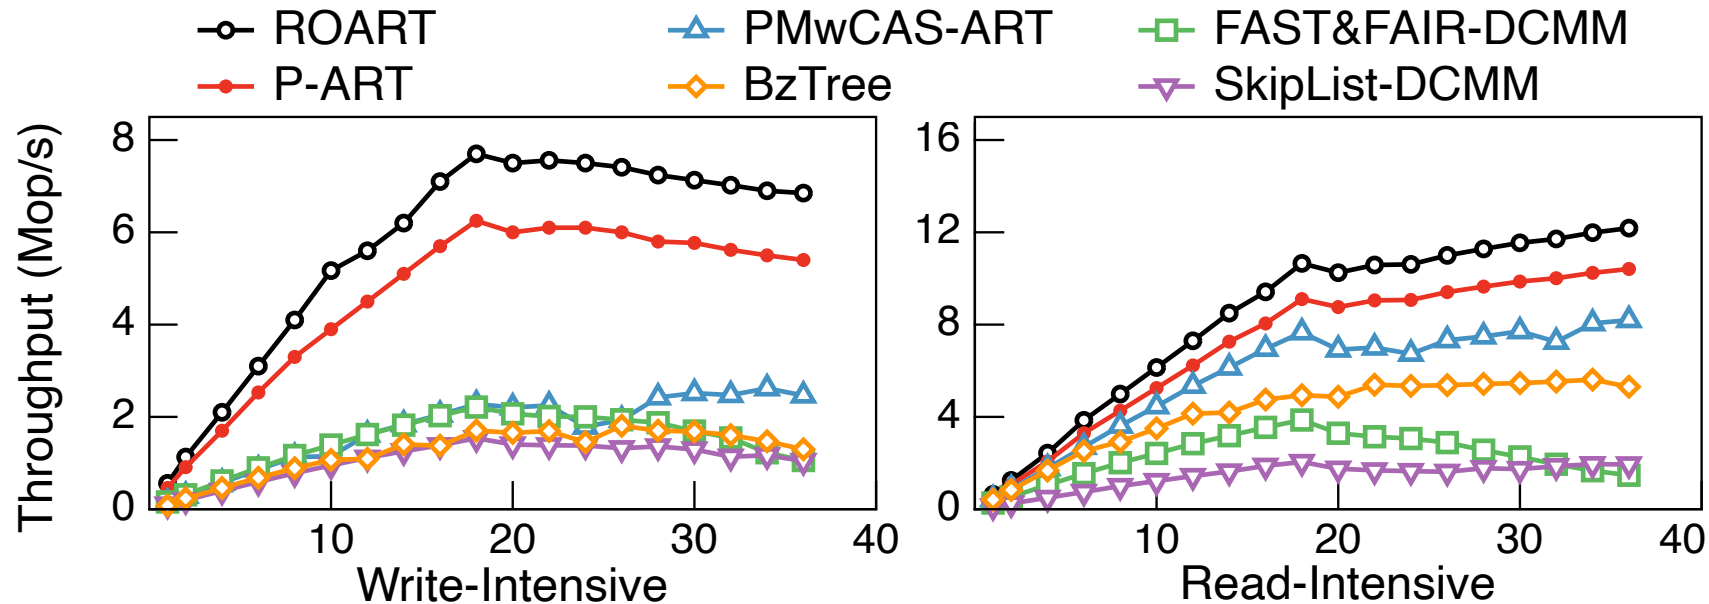
- Higher allocation performance compared with **PMwCAS (PMDK)**
- Less persistent related

- **ROART** achieves **20%-24%** throughput compared with **P-ART** and **PMwCAS-ART** because of leaf compaction mainly → Lower the height of tree and benefit traversal
- **BzTree** has slotted-page node layout → Good cache locality for variable-sized KVs && Binary search

Overall Performance

➤ Methodology

- YCSB benchmark

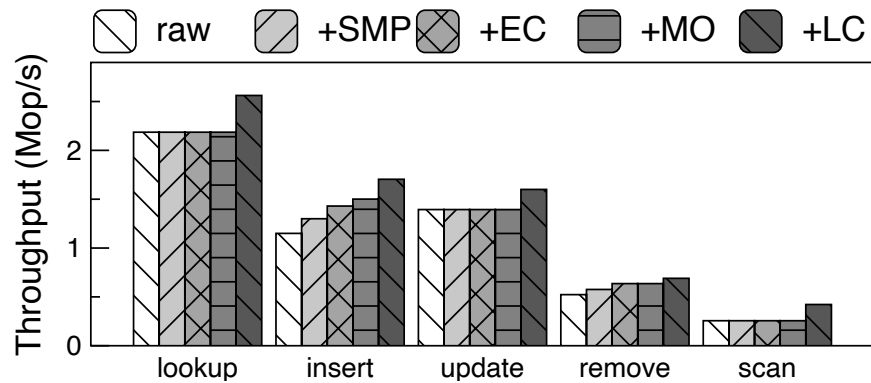


➤ **ROART** achieves **2.78-6.57X** using 36 threads → Less traversal and persistence


Microbenchmark

➤ Each optimization

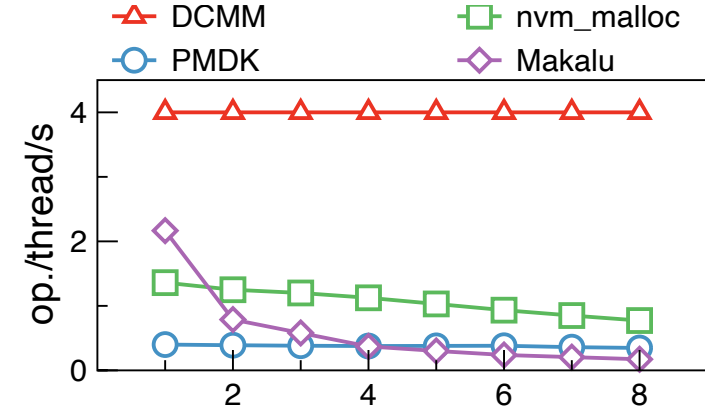
- **SMP**: Selective Metadata Persistence & **EC**: Entry Compression
- **MO**: Minimally Ordered split & **LC**: Leaf Compaction



Performance improvement of each optimization



Range queries with different key numbers



Performance with different allocators


➤ **ROART** with **LC** outperforms the version without **LC** by **1.07~2.01X**

➤ **DCMM**: larger pages in a lock-free manner, page size is only 4K in **Makalu** (lock-based)

Real-world Application Performance

➤ Test in Memcached

- Replace its hash index to three persistent indexes: **ROART**, **P-ART**, **FAST&FAIR**
- Single thread



- **ROART** outperforms **P-ART** and **FAST&FAIR** by up to **1.07X** and **1.38X** in set operations, **1.06X** and **1.19X** in get operations

Conclusion

- **ROART**, a **R**ange **O**ptimized based on **A**daptative **R**adix **T**ree
 - Leaf compaction
 - Entry Compression
 - Selective Metadata Persistence
 - Minimally Ordered Split
 - Delayed Check Memory Management(**DCMM**)