

An Empirical Study of Memory Sharing in Virtual Machines

Sean Barker*, Timothy Wood^,

Prashant Shenoy*, and Ramesh Sitharaman*

*University of Massachusetts Amherst

^The George Washington University

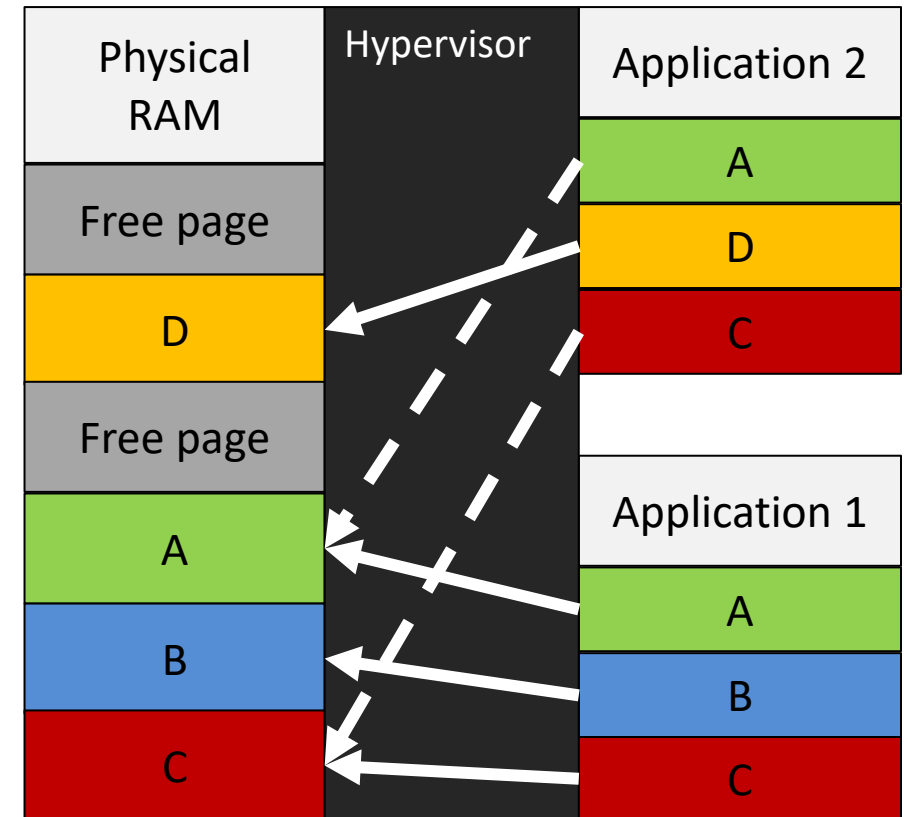
Virtualization and Page Sharing

➤ Virtualization

- Flexible mapping of resources to users
- Efficient resource sharing to maximizing the benefits

➤ Page Sharing

- Eliminate identical pages of memory
- Virtual pages mapped to physical pages
- Hypervisor detects duplicates
- Replaced with copy-on-write references



Content Based Page Sharing Systems

➤ VMware ESX Server [SIGOPS 02]

- **Periodic memory scanning** to detect duplicates
- >30% memory savings

➤ Difference Engine [OSDI 08]

- **Sub-page** sharing and patching
- >60% memory savings

➤ Satori [USENIX ATC 09]

- Sharing of **short-lived** pages
- >90% of possible sharing captured

Problems in Page Sharing

- What levels of sharing are possible in typical **real-world machines**?
- What are the **factors** that impact sharing potential?
 - OS family? Versions? Applications?
- How emerge in technologies impact sharing?
 - New OS technologies?
 - VDI (Virtual desktop infrastructure) farms?
 - LAMP(Linux, **A**pache, **M**ySQL, **P**HP/**P**erl/**P**ython) clusters?
- **Target: Provide practical insights into these questions through a careful study of memory data**

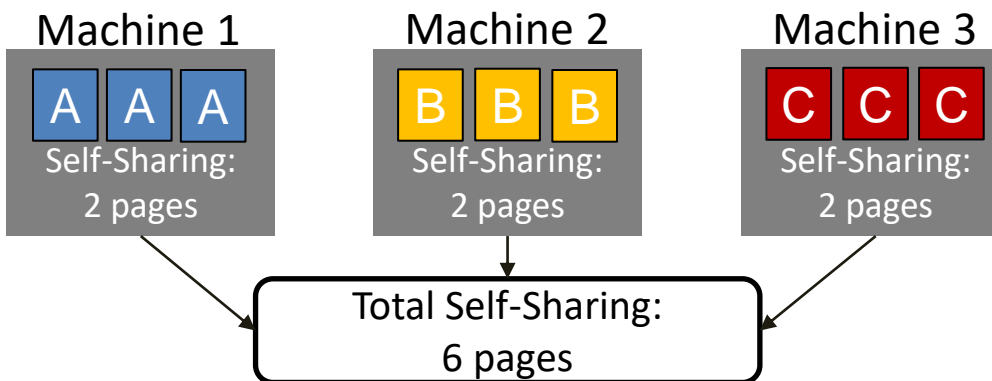
Data Collection

- Traces from uncontrolled machines
 - ~50 real machines (server/desktop mix)
 - Uncontrolled user workloads
 - Memory snapshots every 30 minutes
- Traces from controlled VMs
 - Mixed 32/64bit version of Mac/Win/Linux VMs
 - 3 application setups for each VM:
 - No workload (freshly booted)
 - Server apps (LAMP stack)
 - Desktop apps (office, browser, media player)
- Trace: the content type of page; the process(es) using the page.

Types of Sharing

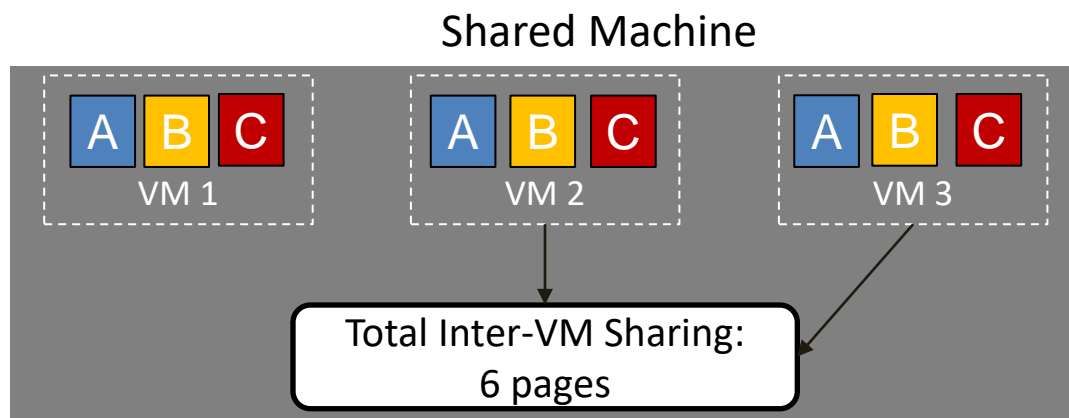
➤ Self-Sharing: sharing **within** individual VMs

- E.g., multiple zero pages



➤ Inter-VM sharing: sharing **across** multiple VMs


- E.g., shared OS states, applications



Real-world Sharing Potential

➤ Self-Sharing:


- Average sharing of 14% in any machine (Excluding zero pages)
- High sharing potential surrounded by long periods of much lower sharing potential



Real-world Sharing Potential

➤ Inter-VM Sharing:

- Observed minimal (<2%) inter-VM sharing potential
- <0.1% sharing in 15 of 21 pairings




Real-world Sharing Potential

- Typical **15%** possible sharing observed
 - Significant, but less than expected from synthetic workloads
- Most (**85+%**) sharing derived from self-sharing
 - What about collocating many VMs?
 - All 7 machines page sharing still **80+%** from self-sharing
- Self-sharing doesn't require virtualization!
 - Could capture it within a VM or nonvirtualized host
- Self-sharing is significant, but what causes it?

Self-Sharing Case Study

➤ Self-Sharing by Process

- **>30%** sharing processes GUI apps/libraries
- **<20%** sharing from other system processes
- Memory footprint likely dominated by GUI




Process self-sharing resulting from user workload

Self-Sharing Case Study

➤ Self-Sharing by Content

- **94%** sharing from libraries and heaps
- Possibly from recreated data structures
- 2.3 MB sharing from single Xorg heap page (~600 copies)




Duplicate data allocations evident in processes

Factors Impacting Sharing

- How do various properties influence sharing?
- Operating system characteristics
 - Family (e.g., Linux or Windows)
 - Version (e.g., Windows XP/7, Ubuntu 10.04/10.10)
 - Architecture (x86 or x64)
- Application setup (LAMP and VDI setups)
- Sharing granularity (number of pages per chunk)
- New OS technologies (e.g., ASLR)

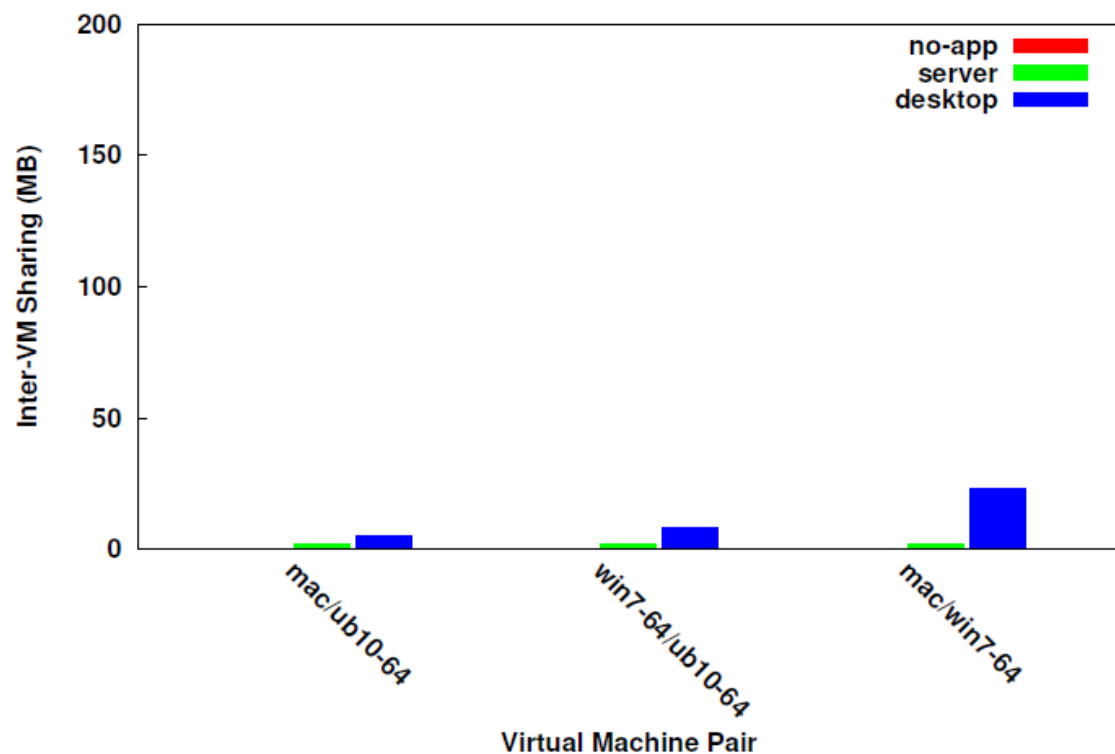
Factors Impacting Sharing

- Self-Sharing across VMs
 - **~100 MB** differences between OS families, major versions (XP/7)
 - **<20 MB** differences between minor versions, architectures

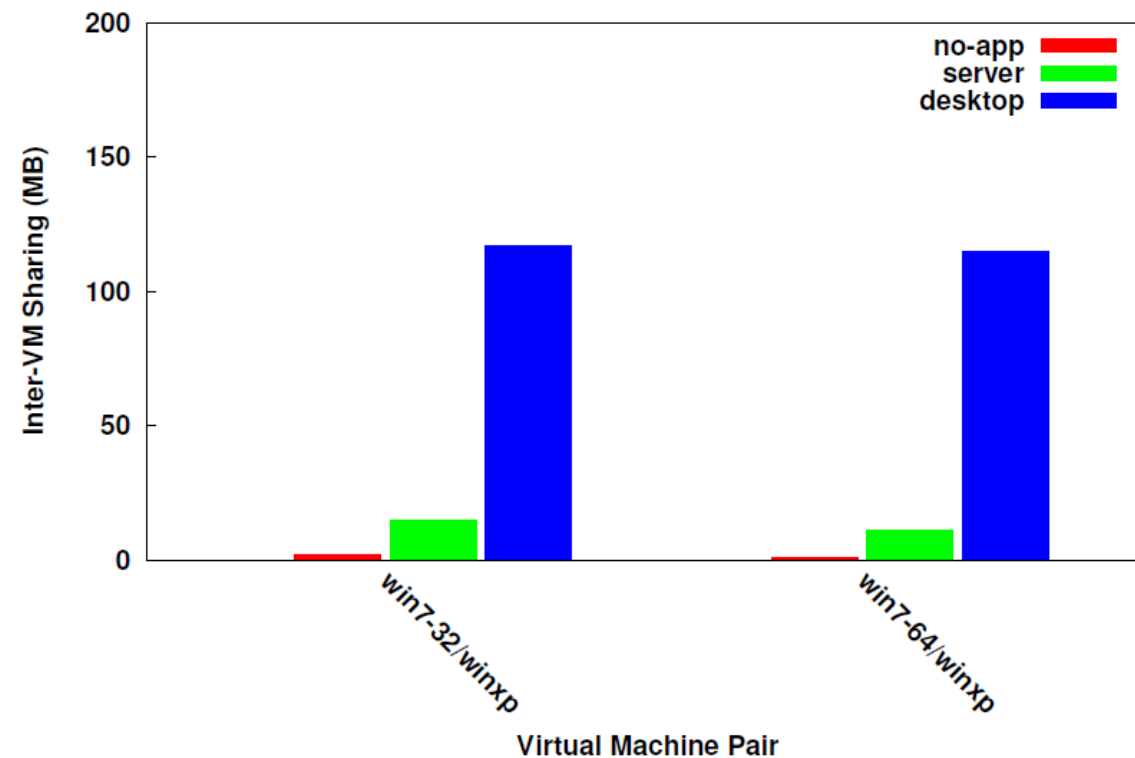


Large self-sharing variations between 'base' OSes

Sharing Across VMs



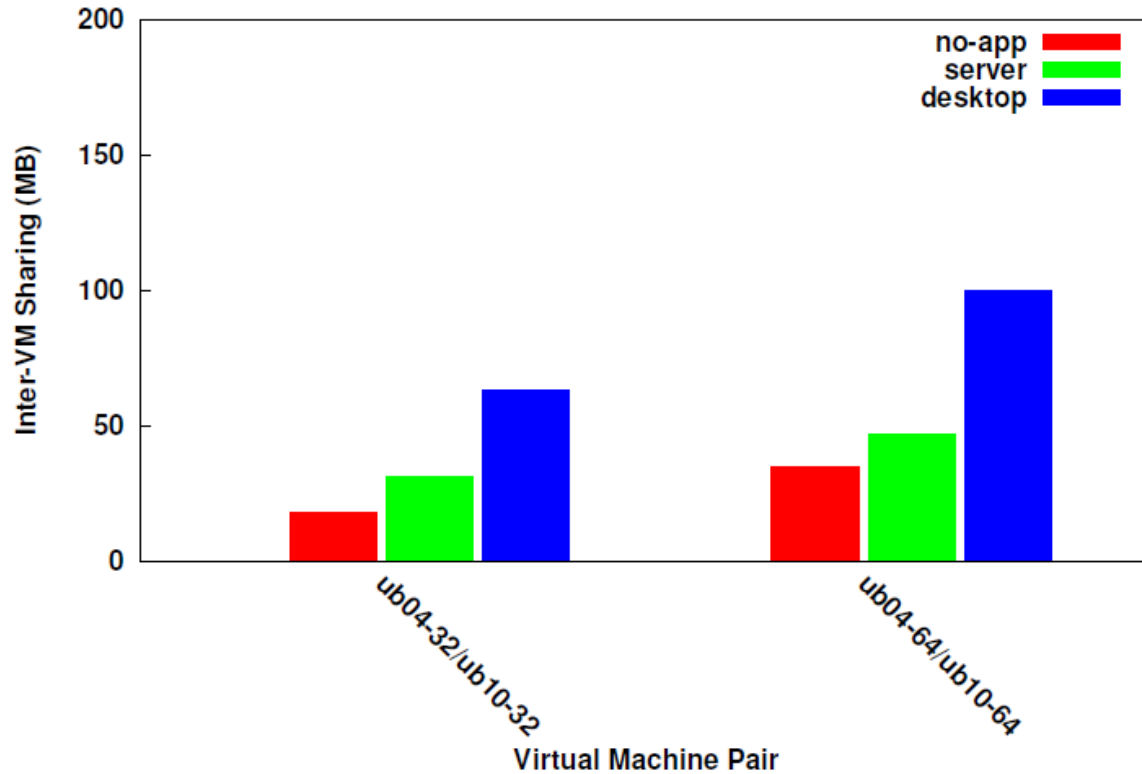
<5 MB across OS families



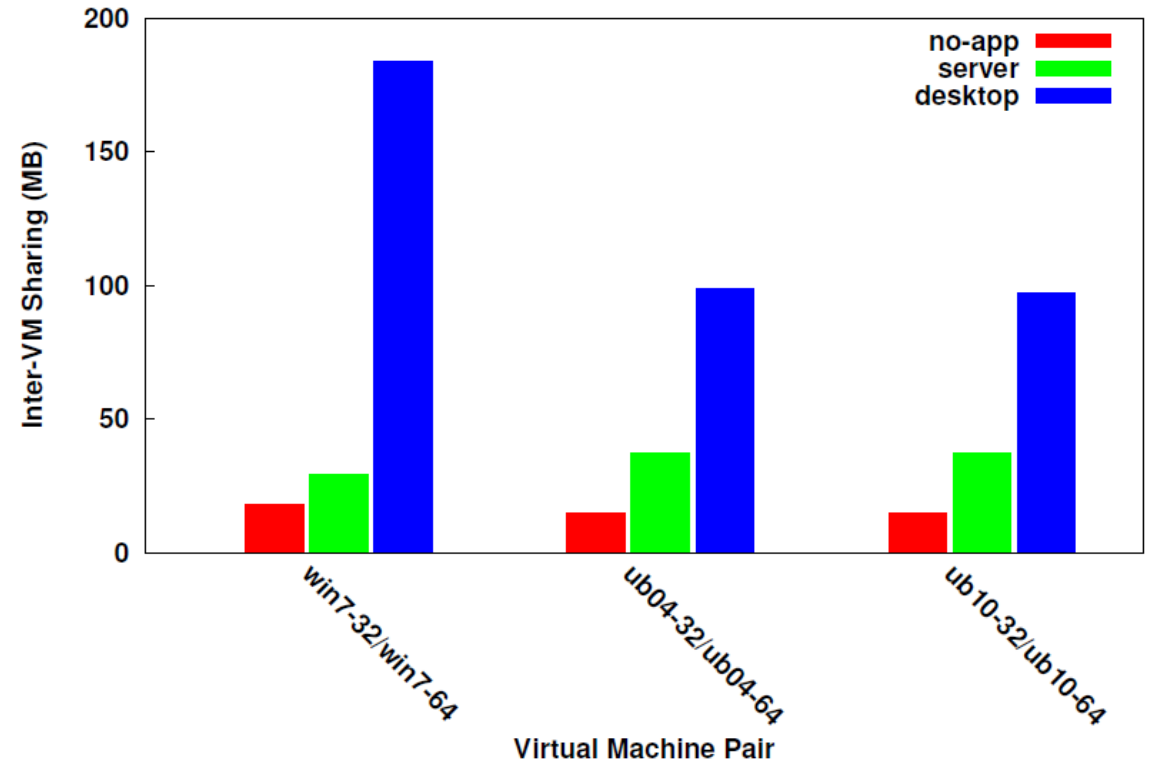
<5 MB base

50+ MB app sharing across major versions

Sharing Across VMs



50+ MB across minor versions



100+ MB across architectures

Hierarchy: family, applications, version, architecture

Sharing Granularity


- Share memory chunks of size k (0.4~2.4) pages
 - Only even page divisions provide decent returns
 - Diminishing benefits from smaller chunk sizes



Tradeoff between overhead and sharing potential

Address Space Layout Randomization

- ASLR scrambles memory to improve system security
 - libraries, code, stack, heap, ...



- Does ASLR have a negative impact on memory sharing?


Address Space Layout Randomization

➤ Impact of 4 ASLR implementations:

- Linux: mainline (2.6.32) and PaX
- Windows 7 (SP1)
- Mac OS X (Lion)

➤ Sharing impact of ASLR

- >10% reduction in three of four cases
- 'Better' (PaX) sharing in Linux worsens impact



ASLR doesn't prevent sharing but does reduce it

Sharing Factor Observations

- **Hierarchy** with respect to sharing potential
 - OS family, application setup, OS version, OS architecture
- **Platform homogeneity**
 - Minimal sharing across heterogeneous systems
 - Significant gains in homogeneous deployments (but still modest absolute levels)
- **Finer-grained** sharing may be leveraged to improve sharing potential
- OS improvements like **ASLR** may reduce sharing

Conclusions

- Study into practical issues of page sharing
 - Examined real-world machines and specific sharing scenarios
- Observed real-world sharing around **15%**
 - Significant, but less than expected
 - Largely self-sharing, for which **no virtualization needed**
- Studied a variety of factors impacting sharing
 - Key role of platform **homogeneity**
 - Varying impact of modifying OS characteristics and applications
 - New technologies may change the impact of sharing