

Improving Restore Speed for Backup Systems that Use Inline Chunk-Based Deduplication

Mark Lillibridge, Kave Eshghi, and Deepavali Bhagwat

HP, Inc

FAST 13

The Problem

➤ Slow restore performance due to **chunk fragmentation**

- Store chunks according to the order of appearance of unique chunks
- Restore the **newer** snapshot, the more random reads, the lower the performance

Snapshot 1:

A	B	C	D	E
---	---	---	---	---

Snapshot 2:

A	B	F	D	E	G
---	---	---	---	---	---

Snapshot 3:

J	A	B	F	D	H	G	I
---	---	---	---	---	---	---	---

Storage :

A	B	C	D	E	F	G	H	I	J
---	---	---	---	---	---	---	---	---	---

The Problem

- Solve the problem via defragment data periodically?
- No layout makes most backups restore fast
 - Backups disagree about optimal order of chunks (MFDedup^[FAST'21])
 - Could focus on last week's backups (RevDedup^[APSYS'13])
- Rearranging chunks is expensive
 - Example: keep last snapshot's chunks in order
- This work: improve speed without rearranging data

Measuring Fragmentation & Restore Speed

- Deduplication storage unit: containers
 - Standard size: 4 MB
 - Observation: reading containers is dominant restore cost (over 80%)
- Measure method: Container read per MB restored
 - Total # of container reads that occur / restored backup size
 - Depends on caching used
 - Uncaptured restore-time variance:
 - Container size is difference
 - File system fragmentation, seek time variance, etc.

Simulating Restoration

➤ Baseline algorithm:

For chunk c in backup B :

 Read in c 's container C

 Extract c 's content from C via offset and size

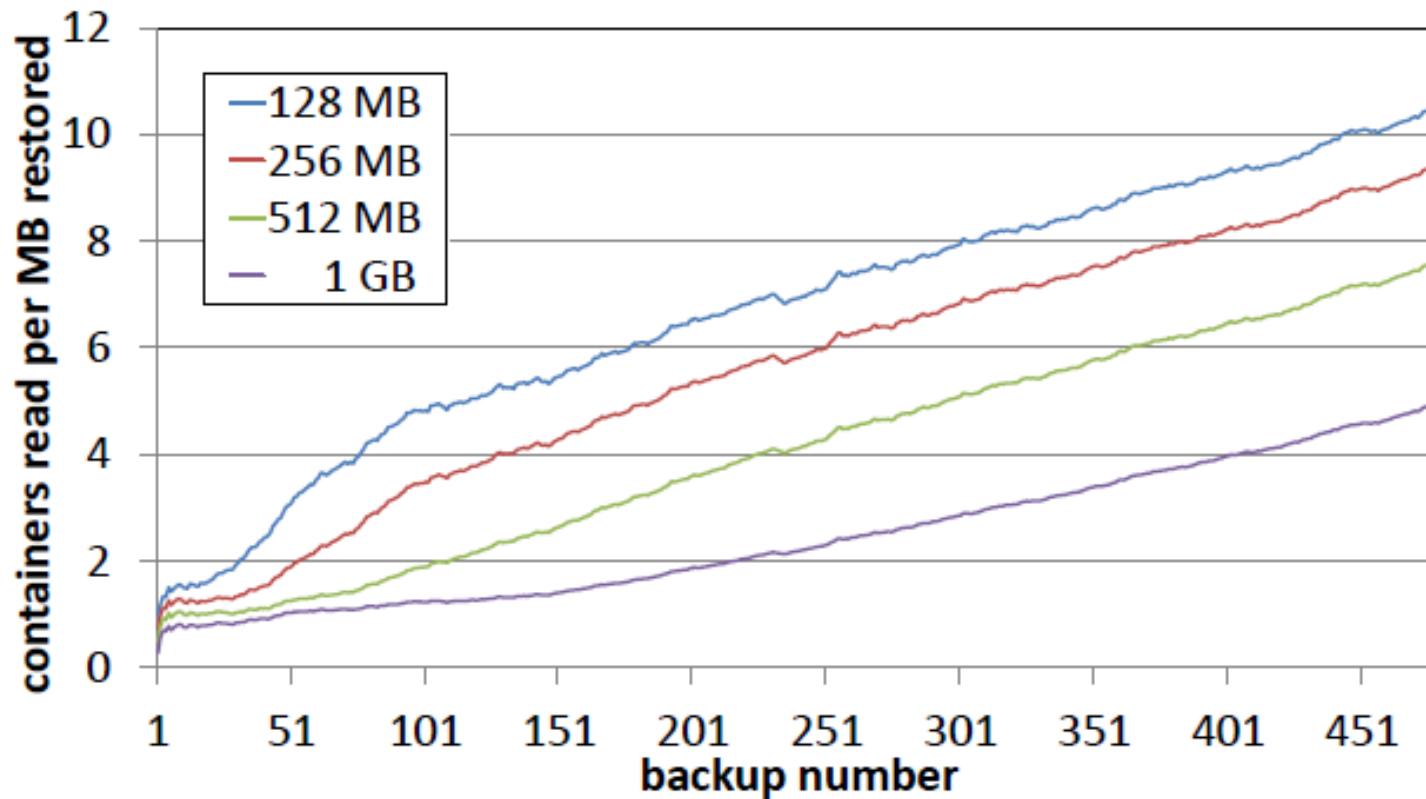
 Send out c 's data

➤ Subtleties:

- LRU caching (Impact of cache size)
- Read entire containers: Read speed \gg Seek time

Fragmentation Over Time

- The curves are different LRU cache sizes



Large caches smooth out small-scale fragmentation

Improving Restore-time Caching

- A new caching method:
 - The forward assembly area method
- Exploits:
 - Perfect knowledge of future accesses
 - Courtesy of the recipe
 - Keeps only needed chunks
- Designed to minimize memory overhead
 - Lots of small variable-sized objects
 - Reduce memory copies


Forward Assembly Area Method

Storage:

Recipe 

Containers   

RAM: 
Container buffer (4MB)


Forward Assembly Area (100MB)



Recipe buffer (~2MB, part of recipe for assembly area)


Forward Assembly Area Method


Storage:

Recipe 

Containers   

RAM: 
Container buffer (4MB)


Recipe buffer (~2MB)


Forward Assembly Area (100MB)

Forward Assembly Area Method

Storage:

Recipe 

Containers   

RAM:



Container buffer (4MB)



Recipe buffer (~2MB)



Forward Assembly Area (100MB)

Forward Assembly Area Method

Storage:

Recipe 

Containers   

RAM: 

Container buffer (4MB)



Recipe buffer (~2MB)



Forward Assembly Area (100MB)

Forward Assembly Area Method

Storage:

Recipe 

Containers   

RAM: 

Container buffer (4MB)



Recipe buffer (~2MB)



Forward Assembly Area (100MB)

Forward Assembly Area Method

Storage:

Recipe 

Containers   

RAM: 

Container buffer (4MB)



Recipe buffer (~2MB)



Forward Assembly Area (100MB)

Forward Assembly Area Method

- Cache only needed chunks
 - Load each container once per slice (100 MB)
 - Uses memory more efficiently
- Moves chunks directly into position
 - No intermediate storage or extra copies
- Better rolling approach
 - A ring-buffer based approach,
 - Send out the continuous filled-in part at the (logical) start of the forward assembly area and then rotate the ring buffer.
 - Not need to reloads around slice boundaries

Forward Assembly Area Rolling Method

Storage:

Recipe 

Containers   

RAM: 

Container buffer (4MB)



Recipe buffer (~2MB)




Forward Assembly Area Rolling Method


Storage:

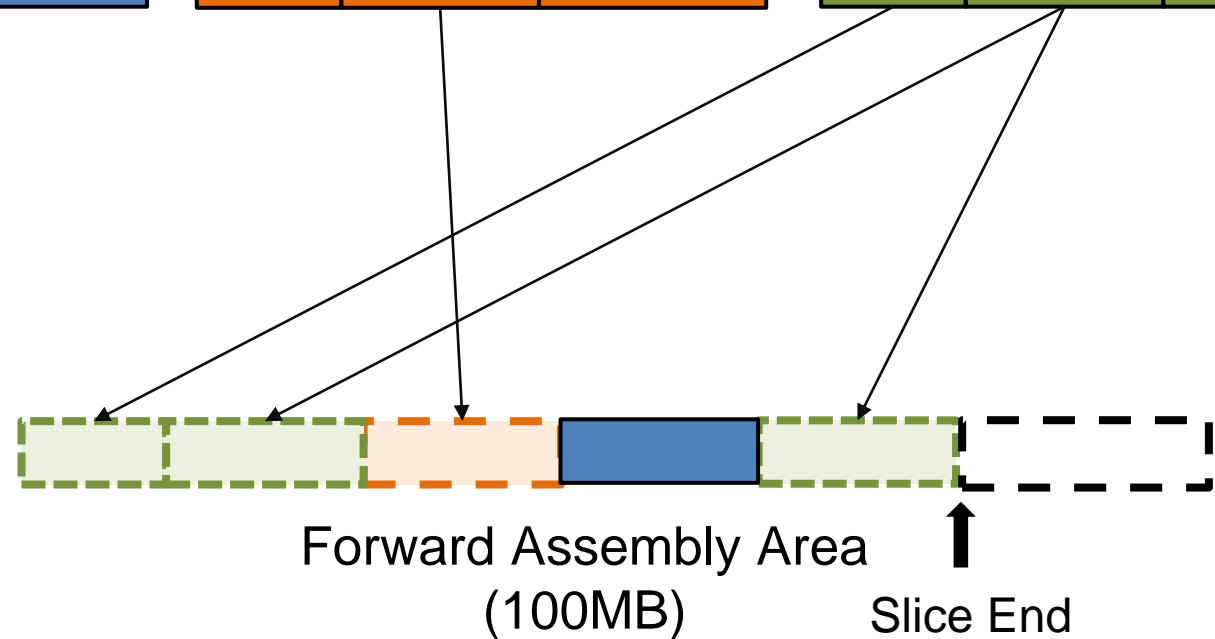
Recipe 

Containers   

RAM:


Container buffer (4MB)


Recipe buffer (~2MB)




Forward Assembly Area Rolling Method

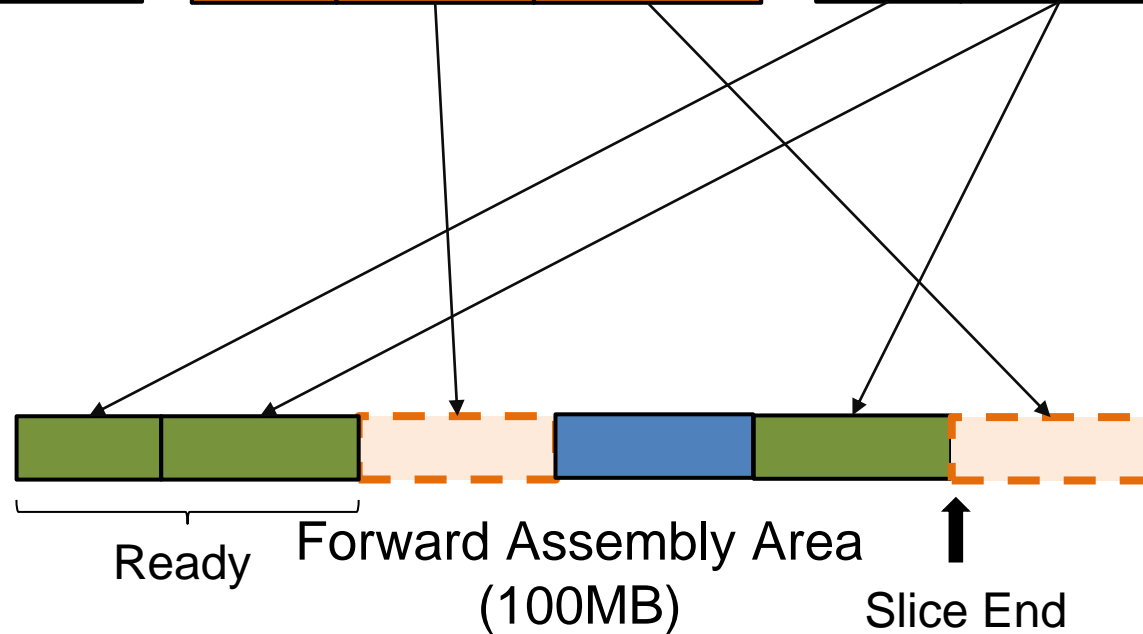
Storage:

Recipe 

Containers   

RAM: 
Container buffer (4MB)


Recipe buffer (~2MB)



Forward Assembly Area Rolling Method

Storage:

Recipe 


Containers   

RAM: 

Container buffer (4MB)



Recipe buffer (~2MB)




Limiting Fragmentation by Reducing Deduplication

- Goal: bound containers read per MB ingested
 - Guarantees minimal restore speed
- The capping method
 - Break input into 20 MB segments
 - Deduplicate segment against at most T containers
 - Limits fragmentation to $(T+5)/20\text{MB}$
 - Here we use 4 MB container, $5 = 20\text{ MB}/4\text{MB}$ (That is, the effective $T=0$ situation)
 - Minimize chunks duplicating by using best T containers
 - Select T according to the number of duplicate chunks contained in each container

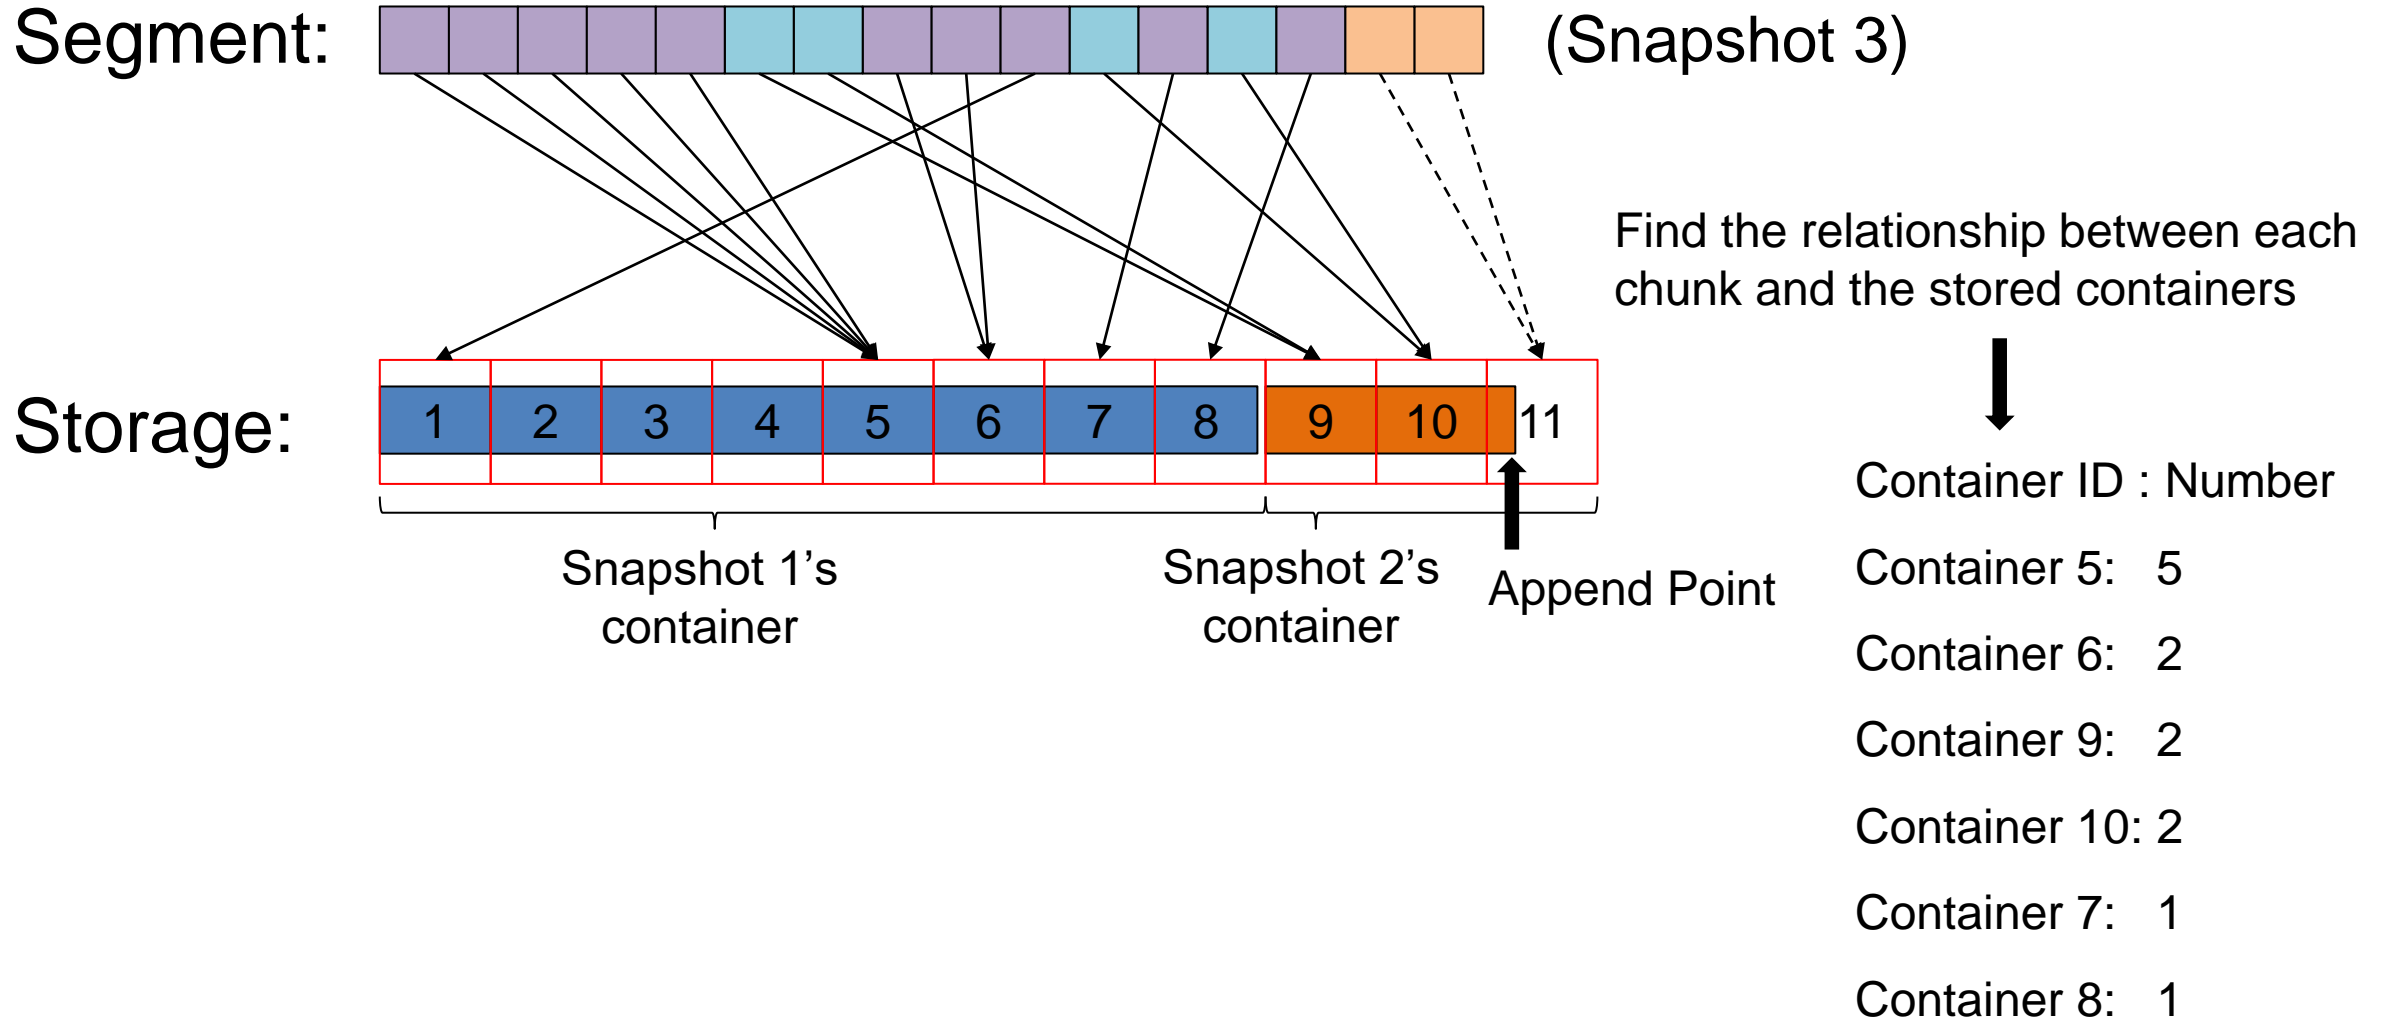
Container Capping

Segment:  (Snapshot 3)

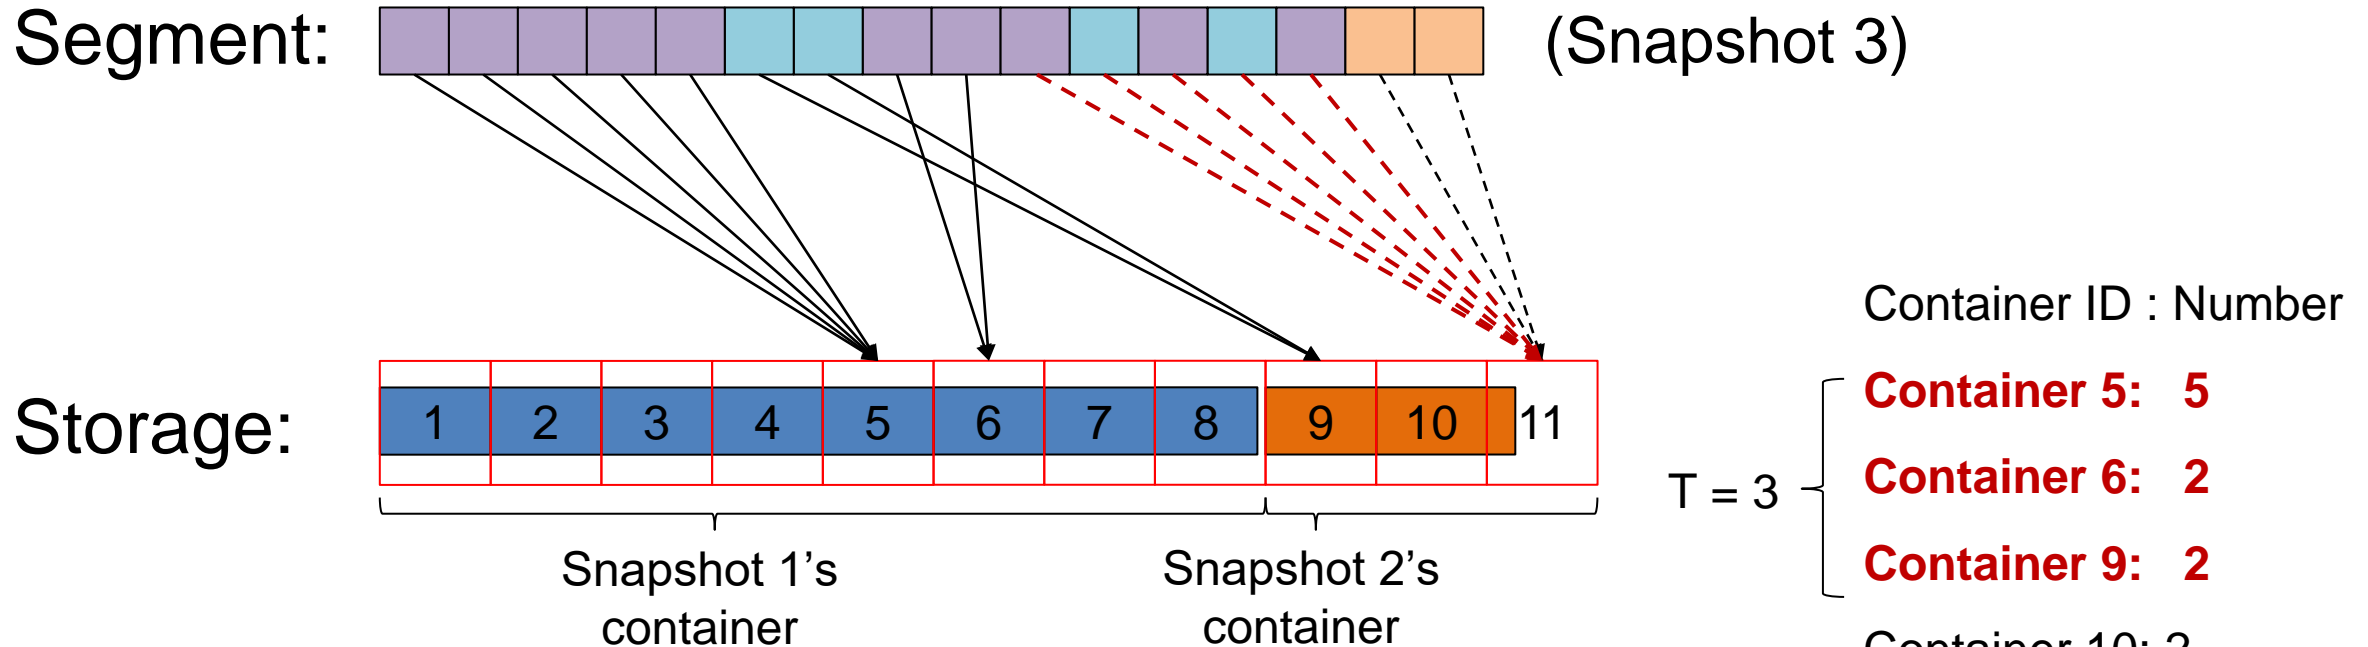
Storage:



Container Capping



Container Capping



➤ Result: 4 containers

- Extra copy: 4 chunks
- No capping: restore need to access **8** containers

Evaluation

➤ Glossary:

- Speed factor: $1 / \text{container read per MB}$
- Deduplication factor: $\text{logic size} / \text{stored size}$



➤ Dataset:

- 2year 212 TB
 - Created based on HP customer data (10 GB snapshot), modify 2% of files by overwriting 10%, and add 200 MB data for each day.
 - 1 full + 4 incrementals per week for 2 years (480 backups)
 - ~4KB mean-size chunks
- Workgroup 3.8 TB
 - Backups of 20 desktop PCs for 91 days.

Evaluation

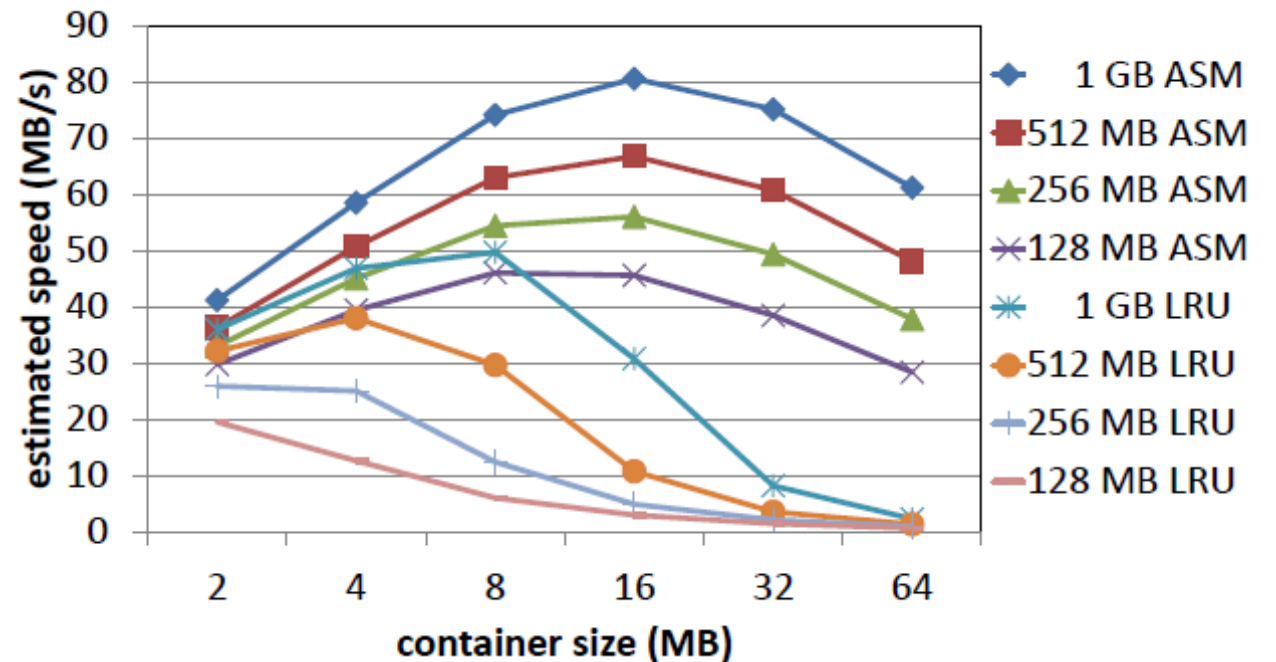
➤ Assembly vs. LRU

- The bigger the speed factor, the better



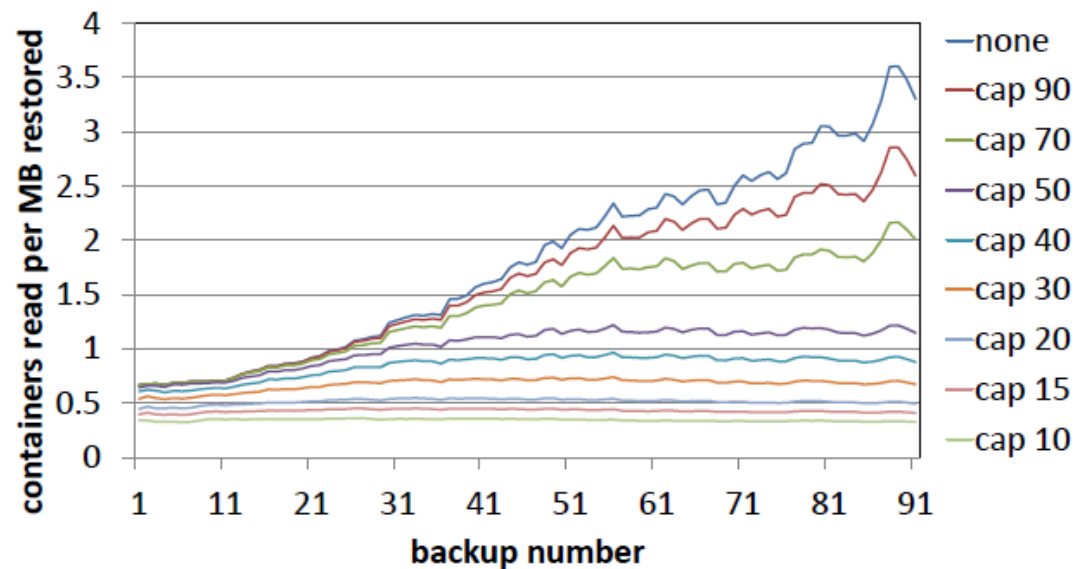
Evaluation

- Estimated speed for a sample system as container size is varied
 - Workgroup-cycle30
 - 4 cache sizes and assembly (rolling) and caching
 - Assume system reads at 1000 MB/s and opens a container in 20 ms
 - All containers were assumed to be full.
 - LRU collapses as container size grows

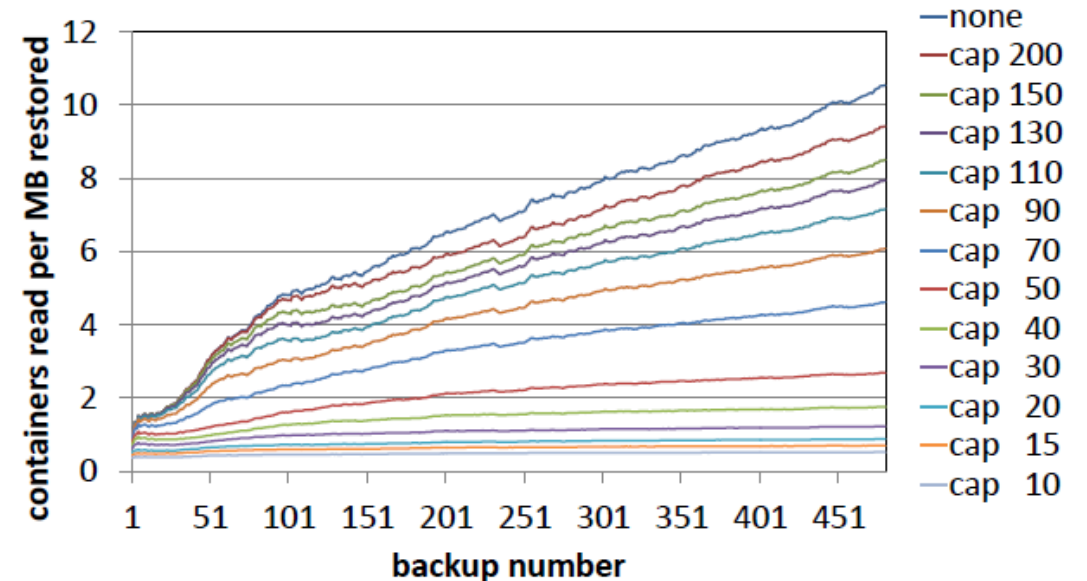


Evaluation

➤ Effect of varying capping level on fragmentation




(a) Workgroup-cycle30



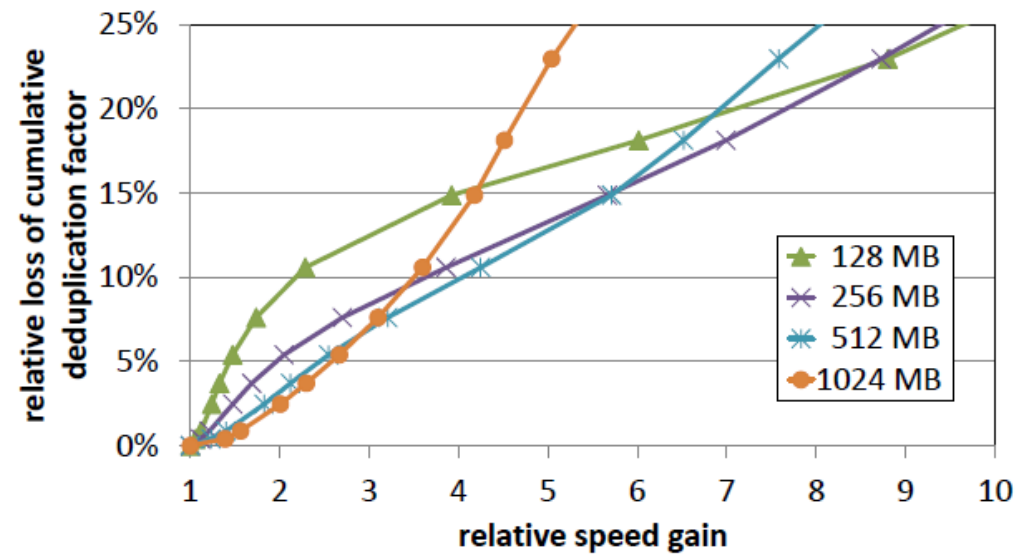
(b) 2year-cycle30

Evaluation

- Relative deduplication loss versus relative speed gain as a result of capping
 - $T = 10, 15, 20, 30, 40, 50, 70, 90, 110, 130, 150, 200, 250, \text{none}$



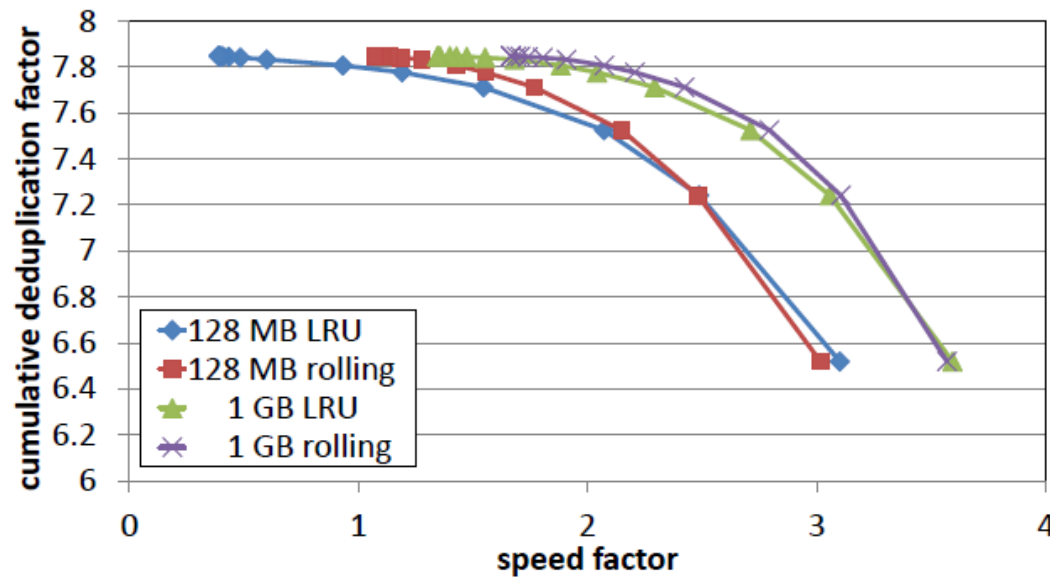
(a) Workgroup-cycle30



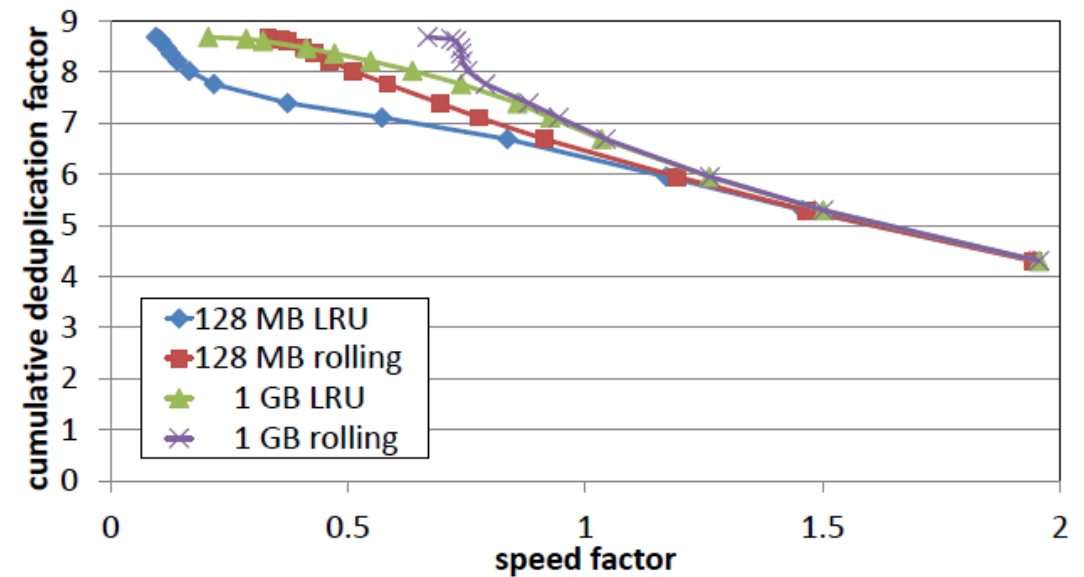
(b) 2year-cycle30

Evaluation

- Comparing deduplication and speed as capping level varies between LRU and assembly
 - $T = 10, 15, 20, 30, 40, 50, 70, 90, 110, 130, 150, 200, 250, \text{none}$



(a) Workgroup-cycle30



(b) 2year-cycle30

Conclusion

- Fragment analysis:
 - Reading containers is dominant restore cost
 - Measure fragment by container read per MB
- Caching: Reduce seeks through better caching
 - Forward assembly area
 - Load each container once per slice (100 MB)
- Capping: Limiting fragmentation by reducing deduplication
 - Deduplicate segment against at most T containers
 - Bound containers read per MB ingested